





Regular Expressions Overview



What is a Regular Expression?

A regular expression defines a pattern of characters.

Can be used for:

Validation



Data Extraction

Data Cleaning




One pattern can match one or many sets of characters

English	Pattern	Matches	Does Not Match
4 numbers in a row	<code>\d\d\d\d</code> or <code>\d{4}</code>	1234 2222 3333	a1234 AAsaaaa 123
2 numbers, a slash, two numbers, a slash, 4 numbers	<code>\d\d\d\d\d\d\d</code> or <code>\d{2}\d{2}\d{4}</code>	11/01/2013 10/22/2015 23/45/2222	11/1/2013 1/11/2015 aa/aa/aaaa dsifjosdijfoas



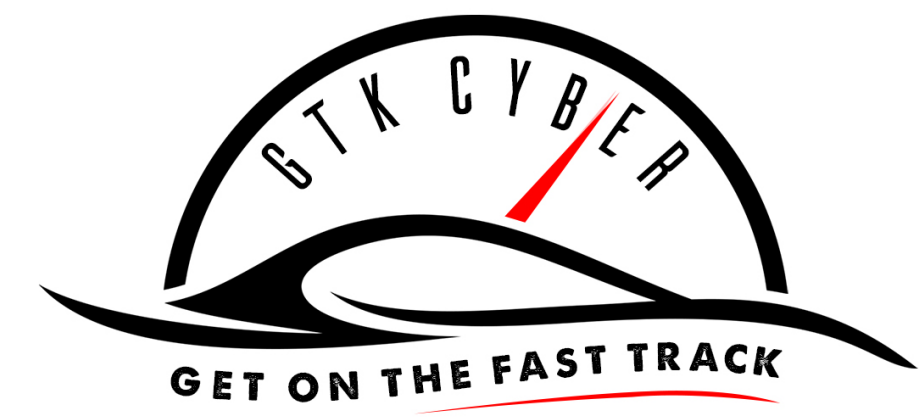
regexpal.com



Challenge 1

Let's write a pattern that matches a date. Such as...

07/30/2016



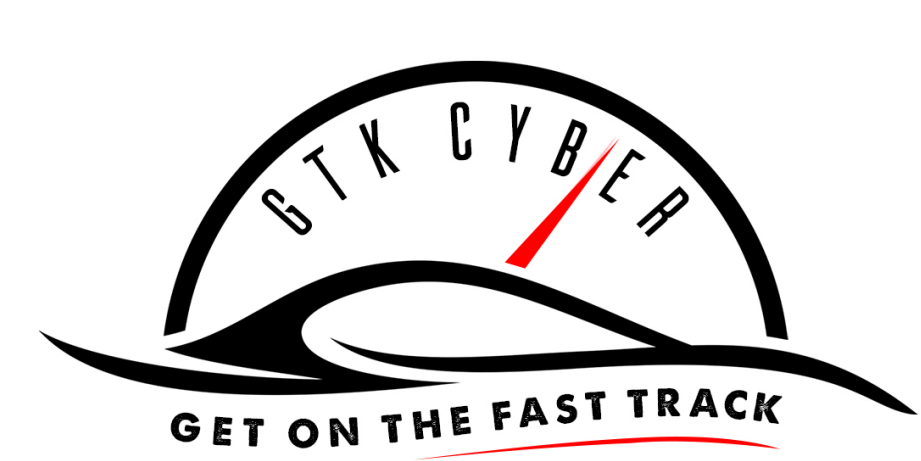
Each Regex Character Represents a Character in a String



10 boxes for 10 characters



But this will only
match our one date.



Literal vs Special Characters

Literal Characters

0 7 / 3 0 / 2 0 1 6

Character Sets

\d \d / \d \d / \d \d \d \d

Wildcards

. . / . . /



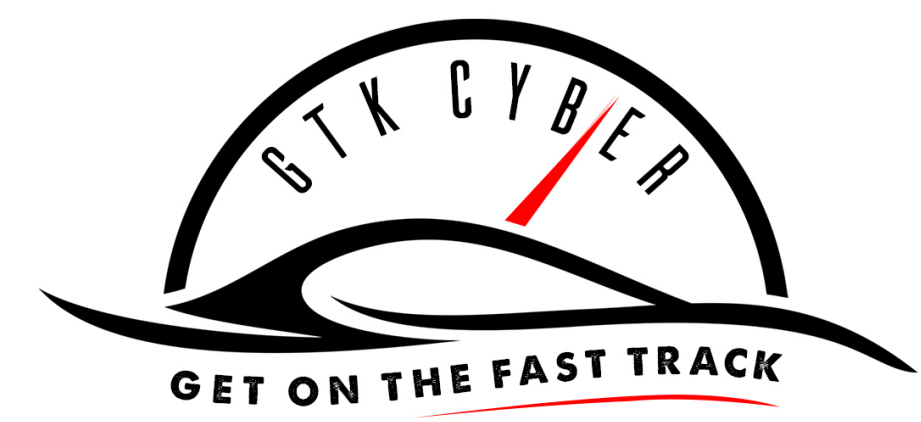
Literal Characters

- Escape certain characters that have special meaning
 - \ can define a character set or escape a special character (\d or \. or \\)



Character Sets

- Can explicitly define a set of characters
 - [aeiou]
- Can define a range of characters
 - [a-z0-9]
- Can represent a set of characters
 - \d
 - \w
- Can represent *not* characters
 - [^aeiou]
 - \D



Shorthand for Character Sets

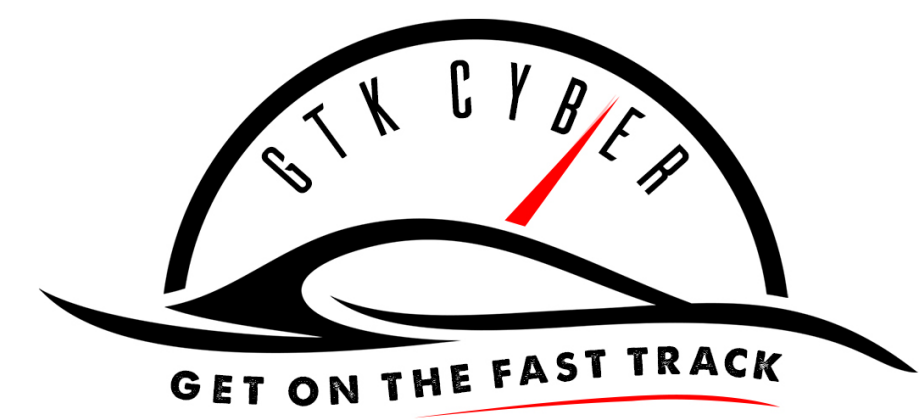
There are shortcuts for commonly used character sets:

Shortcut	Definition	Example
<code>\s</code>	Any whitespace character	<code>/a\s b/</code> matches: <code>a b</code>
<code>\S</code>	Any non-whitespace character	<code>/a\S b/</code> matches : <code>abb</code>
<code>\d</code>	Any digit	<code>\d\d-\d</code> matches <code>12-3</code>
<code>\D</code>	Any non-digit	<code>/a\Db/</code> matches <code>aBc</code> or <code>abc</code>
<code>\w</code>	Any alpha-numeric character	
<code>\W</code>	Any non-alpha-numeric character	




Wildcard

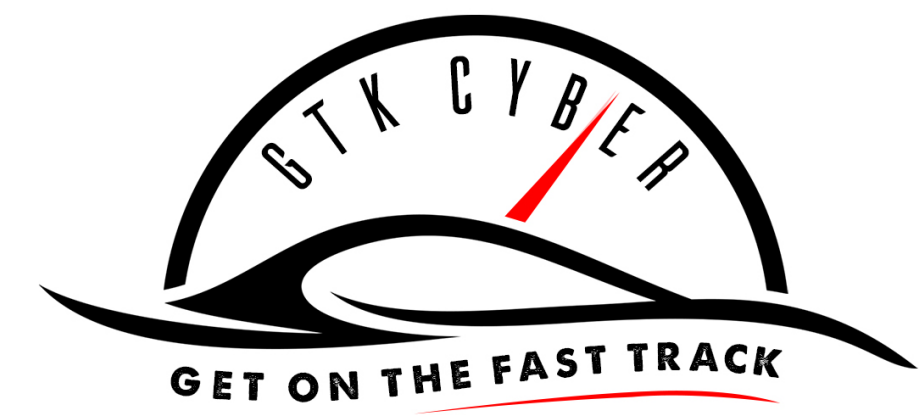




Repetition


We don't have to use 10 boxes when we have repeated characters.





Grouping

Parentheses articulate groups of characters that can be extracted or repeated.





Challenge 2

Let's write a pattern that matches an email. Such as...

`guy9@gmail.com`

You try! Write a pattern that uses characters sets and repetition to match the email.



Less defined repetition

Literal Characters

g	u	y	9	@	g	m	a	i	l	\.	c	o	m
---	---	---	---	---	---	---	---	---	---	----	---	---	---

Character Sets

\w	\w	\w	\w	@	\w	\w	\w	\w	\w	\.	\w	\w	\w
----	----	----	----	---	----	----	----	----	----	----	----	----	----

Repetition

\w {4}	@	\w {5}	\.	\w {3}
--------	---	--------	----	--------

But what if there are 6 characters in the first part of the email?



Question, Star, and Plus

? match the previous character 0 or 1 times

***** match the previous character 0 or more times

+ match the previous character 1 or more times



Greedy vs Lazy

Sometimes `.+` can match too much.

If we throw `<.+>` at `<h1>Welcome</h1>` to find opening tags, we get back the entire string when we only wanted the beginning.

Using `.+?` makes the `+` lazy, meaning it will only grab as many characters are needed in order to continue the match.



Application

$\boxed{\backslash w} \{4\} \boxed{@} \boxed{\backslash w} \{5\} \boxed{\.} \boxed{\backslash w} \{3\}$

$\boxed{\backslash w} + \boxed{@} \boxed{\backslash w} + \boxed{\.} \boxed{\backslash w} +$



Exercises

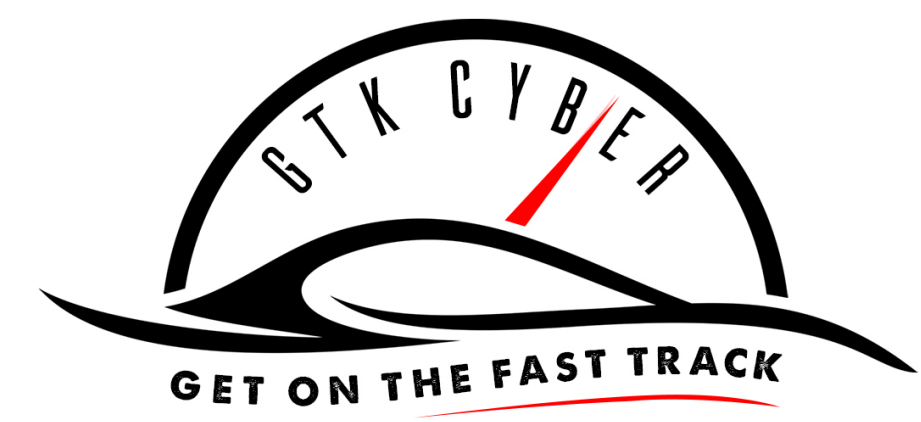
Write regular expressions for the following:

- Filenames in the following format: `yyymmdd-data.xls`
- IP Addresses in the format `XXX.XXX.XXX.XXX`
- Social Security Numbers in the format `XXX-XX-XXXX`
- Any 4 letter word beginning with a vowel
- Any 4 letter word with a number at the end



Regex in Python

- Python has regex support via the re module, which must be imported.
- The re module has four basic functions
 - `match(<pattern>, <text>)`: finds the **first** occurrence of the pattern in the given text.
 - `search(<pattern>, <text>)`: finds any occurrence of the pattern in the given text
 - `findall(<pattern>, <text>)/finditer(<pattern>, <text>)`: finds all occurrences of the pattern in a given text.
 - `split(<pattern>, <text>)`: Splits the text by the regex.
 - `sub(<old>, <new>, <text>)`: Replaces old with the new in the given text.



Regex Option Flags

Flag	Description
re.I / re.IGNORECASE	Performs case insensitive matching
re.L / re.LOCALE	Interprets words according to locale
re.M / re.MULTILINE	Make begin consider each line
re.S / re.DOTALL	Makes a period match any character including a newline.
re.U / re.UNICODE	Interprets letters according to the Unicode character set. This flag affects the behavior of \w, \W, \b, \B.
re.X / re.VERBOSE	Allows comments in regex



Regex in Python

```
import re

text = "some 4444 text"
regex = "\d{4}"

matchObj = re.match(regex, text, re.U)

if matchObj:
    # Successful Match

else:
    #No match
```



Regex in Python

```
import re

text = "some 4444 text"
regex = "\d{4}"

matchObj = re.match(regex, text, re.U)

if matchObj:
    # Successful Match

else:
    #No match
```




Regex in Python

```
import re

text = "some 4444 text"
regex = r"\d{4}"

# Compiling Regex will improve performance
compiled_regex = re.compile(regex)
matchObj = compiled_regex.search(regex, text, re.U)

if matchObj:
    # Successful Match

else:
    #No match
```



Grouping Parentheses

- When you put parens around sections of a regex you use these to extract parts of the text
- Python uses the `.group(n)` function to access parts of a match
- `group(0)` will get you the entire matched text, whereas `group(1)` gets the first match.



Extracting Data with Regex

```
import re

emailAddress = "account@domain.com"
emailRegex = r"(\w+)@(\w+\.\w+)"

emailMatch = re.search(emailRegex, emailAddress)

if emailMatch:
    account = emailMatch.group(1)
    domain = emailMatch.group(2)
    completeEmail = emailMatch.group(0)
else:
    #No match
```



Back References

- Back references allow you to refer to previously matched blocks of text.
- Python uses the syntax `\1`, `\2`, `\3` in a regex to refer to previously matched groups
- Can be used in `re.sub()` to re-arrange matched parts.



In Class Exercise

Please take 20 minutes and complete
Worksheet 0.1: Regular Expressions in Python



Questions?